

PSTAT131_HW6

Tao Wang

2022-05-20

The goal of this assignment is to build a statistical learning model that can predict the **primary type** of a Pokémon based on its generation, legendary status, and six battle statistics.

Note: Fitting ensemble tree-based models can take a little while to run. Consider running your models outside of the .Rmd, storing the results, and loading them in your .Rmd to minimize time to knit.

Loading Packages

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(tidymodels)

## -- Attaching packages ----- tidymodels 0.2.0 --
## v broom      0.7.12      v rsample      0.1.1
## v dials      0.1.0      v tune         0.2.0
## v infer      1.0.0      v workflows    0.2.6
## v modeldata  0.1.1      v workflowsets 0.2.1
## v parsnip    0.2.1      v yardstick    0.0.9
## v recipes    0.2.0

## -- Conflicts ----- tidymodels_conflicts() --
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed()  masks stringr::fixed()
## x dplyr::lag()      masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step()   masks stats::step()
## * Search for functions across packages at https://www.tidymodels.org/find/

library(ISLR)
library(rpart.plot)

## Loading required package: rpart
```

```

##
## Attaching package: 'rpart'
## The following object is masked from 'package:dials':
##
##      prune
library(vip)

##
## Attaching package: 'vip'
## The following object is masked from 'package:utils':
##
##      vi
library(janitor)

##
## Attaching package: 'janitor'
## The following objects are masked from 'package:stats':
##
##      chisq.test, fisher.test
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##      combine
## The following object is masked from 'package:ggplot2':
##
##      margin
library(xgboost)

##
## Attaching package: 'xgboost'
## The following object is masked from 'package:dplyr':
##
##      slice
library(corrplot)

## corrplot 0.92 loaded
library(ISLR2)

##
## Attaching package: 'ISLR2'
## The following objects are masked from 'package:ISLR':
##
##      Auto, Credit

```

```
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
## Loaded glmnet 4.1-4
library(ranger)

##
## Attaching package: 'ranger'
## The following object is masked from 'package:randomForest':
##
##     importance
tidymodels_prefer()
```

Exercise 1

Read in the data and set things up as in Homework 5:

- Use `clean_names()`
- Filter out the rarer Pokémon types
- Convert `type_1` and `legendary` to factors

Do an initial split of the data; you can choose the percentage for splitting. Stratify on the outcome variable.

Fold the training set using v -fold cross-validation, with $v = 5$. Stratify on the outcome variable.

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`:

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.

```
# read the data
pokemon <- read.csv(file = "data/Pokemon.csv")

# Use its `clean_names()` function on the Pokémon data
pokemon <- clean_names(pokemon)
head(pokemon)
```

```
##   x          name type_1 type_2 total hp attack defense sp_atk sp_def
## 1 1      Bulbasaur Grass Poison  318 45    49    49    65    65
## 2 2      Ivysaur  Grass Poison  405 60    62    63    80    80
## 3 3      Venusaur Grass Poison  525 80    82    83   100   100
## 4 3 VenusaurMega Venusaur Grass Poison  625 80   100   123   122   120
## 5 4      Charmander  Fire      309 39    52    43    60    50
## 6 5      Charmeleon  Fire      405 58    64    58    80    65
##   speed generation legendary
## 1    45          1      False
## 2    60          1      False
## 3    80          1      False
```

```
## 4      80          1      False
## 5      65          1      False
## 6      80          1      False

# Filter out the rarer Pokémon types
pokemon <- pokemon %>%
  filter(type_1 %in% c('Bug', 'Fire', 'Grass', 'Normal', 'Water', 'Psychic'))

# Convert `type_1` and `legendary` to factors
pokemon$type_1 <- factor(pokemon$type_1)
pokemon$legendary <- factor(pokemon$legendary)

# Do an initial split of the data; you can choose the percentage for splitting.
# Stratify on the outcome variable.
set.seed(1234)
pokemon_split <- initial_split(pokemon, prop = 0.80, strata = type_1)
pokemon_train <- training(pokemon_split)
pokemon_test <- testing(pokemon_split)

# Fold the training set using *v*-fold cross-validation, with `v = 5`.
# Stratify on the outcome variable.
set.seed(1234)
pokemon_folds <- vfold_cv(pokemon_train, v = 5, strata = type_1)

# Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`,
# `attack`, `speed`, `defense`, `hp`, and `sp_def`:
# - Dummy-code `legendary` and `generation`;
# - Center and scale all predictors.
pokemon_recipe <- recipe(type_1 ~ legendary + generation + sp_atk +
  attack + speed + defense + hp + sp_def, pokemon_train) %>%
  step_dummy(legendary, generation) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```

Exercise 2

Create a correlation matrix of the training set, using the `corrplot` package. *Note: You can choose how to handle the continuous variables for this plot; justify your decision(s).*


What relationships, if any, do you notice? Do these relationships make sense to you?

```
head(pokemon_train)
```

##	x	name	type_1	type_2	total	hp	attack	defense	sp_atk	sp_def
## 15	11	Metapod	Bug		205	50	20	55	25	25
## 16	12	Butterfree	Bug	Flying	395	60	45	50	90	80
## 17	13	Weedle	Bug	Poison	195	40	35	30	20	20
## 18	14	Kakuna	Bug	Poison	205	45	25	50	25	25
## 19	15	Beedrill	Bug	Poison	395	65	90	40	45	80
## 20	15	BeedrillMega	Bug	Poison	495	65	150	40	15	80
##	speed	generation	legendary							
## 15	30	1	False							
## 16	70	1	False							
## 17	50	1	False							
## 18	35	1	False							
## 19	75	1	False							

```
## 20    145          1    False
# notice we should remove variable x and generation
# since x is the index and generation is categorical
pokemon_train%>%
  select(is.numeric,-x,-generation) %>%
  cor() %>%
  corrplot(type = 'lower', diag = FALSE,
           method = 'color')

## Warning: Predicate functions must be wrapped in `where()`.
##
## # Bad
## data %>% select(is.numeric)
##
## # Good
## data %>% select(where(is.numeric))
##
## i Please update your code.
## This message is displayed once per session.
```



- What relationships, if any, do you notice?
- Answer:
 1. Variable total has positive relationships with variable hp, attack, defense, sp_atk, sp_def and speed.
 2. Variable hp has positive relationships with variable attack, defense, sp_atk and sp_def. Variable hp also has a little positive relationship with variable speed.
 3. Variable attack has positive relationships with variable defense, sp_atk, sp_def and speed.

- 4. Variable defense has positive relationships with variable sp_atk and sp_def. Variable defense also has a little positive relationship with variable speed.
- 5. Variable sp_atk has positive relationships with variable sp_def and speed.
- 6. Variable sp_def has a little positive relationship with variable speed.
- Do these relationships make sense to you?
- Answer:
- Yes, these relationships make sense to me. For example, variable total has positive relationships with variable hp, attack, defense, sp_atk, sp_def and speed. It is because variable total is a general guide to how strong a pokemon is, and how strong a pokemon is depends on variable hp, attack, defense, sp_atk, sp_def and speed. So, it is why the variable total has positive relationships with variable hp, attack, defense, sp_atk, sp_def and speed.

Exercise 3

First, set up a decision tree model and workflow. Tune the `cost_complexity` hyperparameter. Use the same levels we used in Lab 7 – that is, `range = c(-3, -1)`. Specify that the metric we want to optimize is `roc_auc`.

Print an `autoplot()` of the results. What do you observe? Does a single decision tree perform better with a smaller or larger complexity penalty?

```
# set up a decision tree model
tree_spec <- decision_tree() %>%
  set_engine("rpart")

class_tree_spec <- tree_spec %>%
  set_mode("classification") %>%
  set_args(cost_complexity = tune()) # Tune the `cost_complexity` hyperparameter


# set up a decision tree workflow.
class_tree_wf <- workflow() %>%
  add_model(class_tree_spec) %>%
  add_recipe(pokemon_recipe)

# Use the same levels we used in Lab 7 -- that is, `range = c(-3, -1)`.
param_grid <- grid_regular(cost_complexity(range = c(-3, -1)), levels = 10)

# Specify that the metric we want to optimize is `roc_auc`
tune_res <- tune_grid(
  class_tree_wf,
  resamples = pokemon_folds,
  grid = param_grid,
  metrics = metric_set(roc_auc)
)
```

```
## ! Fold1: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold2: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold3: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold4: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold5: preprocessor 1/1: The following variables are not factor vectors and wil...
```

```
# Print an `autoplot()` of the results.
autoplot(tune_res)
```



What do you observe? * Answer: * We can see that as the value of the cost-complexity parameter increases, the ROC_AUC value will finally decrease.

- To be more specific,
- when the cost-complexity parameter less than 0.00464, as the the cost-complexity parameter increases, the ROC_AUC value increases (a little bit);
- when the cost-complexity parameter is around 0.00464 to 0.0129, as the the cost-complexity parameter increases, the ROC_AUC value decreases (a little bit);
- when the cost-complexity parameter is around 0.0129 to 0.0599, as the the cost-complexity parameter increases, the ROC_AUC value increases (a little bit)
- when the cost-complexity parameter is greater than 0.0599, as the the cost-complexity parameter increases, the ROC_AUC value decreases (very fast).
- Does a single decision tree perform better with a smaller or larger complexity penalty?
- Answer:
- According to the graph, we know that a single decision tree perform better with a smaller complexity penalty.

Exercise 4

What is the `roc_auc` of your best-performing pruned decision tree on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```
# What is the `roc_auc` of your best-performing boosted tree model on the folds?
arrange(collect_metrics(tune_res), desc(mean))
```

```
## # A tibble: 10 x 7
##   cost_complexity .metric .estimator mean      n std_err .config
##   <dbl> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1      0.00464 roc_auc hand_till 0.644     5 0.0109 Preprocessor1_Model04
## 2      0.00774 roc_auc hand_till 0.639     5 0.0124 Preprocessor1_Model05
## 3      0.00278 roc_auc hand_till 0.639     5 0.00891 Preprocessor1_Model03
## 4      0.001   roc_auc hand_till 0.638     5 0.00916 Preprocessor1_Model01
## 5      0.00167 roc_auc hand_till 0.638     5 0.00916 Preprocessor1_Model02
## 6      0.0359  roc_auc hand_till 0.628     5 0.0139 Preprocessor1_Model08
## 7      0.0215  roc_auc hand_till 0.625     5 0.0141 Preprocessor1_Model07
## 8      0.0599  roc_auc hand_till 0.623     5 0.0100 Preprocessor1_Model09
## 9      0.0129  roc_auc hand_till 0.610     5 0.0124 Preprocessor1_Model06
## 10     0.1     roc_auc hand_till 0.5       5 0       Preprocessor1_Model10
```

- What is the roc_auc of your best-performing pruned decision tree on the folds?
- Answer: The roc_auc of my best-performing pruned decision tree on the folds is 0.644.

Exercise 5

Using `rpart.plot`, fit and visualize your best-performing pruned decision tree with the *training* set.

```
# according to the professor, we can just use select_best() method
best_complexity <- select_best(tune_res)
```

```
class_tree_final <- finalize_workflow(class_tree_wf, best_complexity)
```

```
class_tree_final_fit <- fit(class_tree_final, data = pokemon_train)
```

```
## Warning: The following variables are not factor vectors and will be ignored:
## `generation`
```

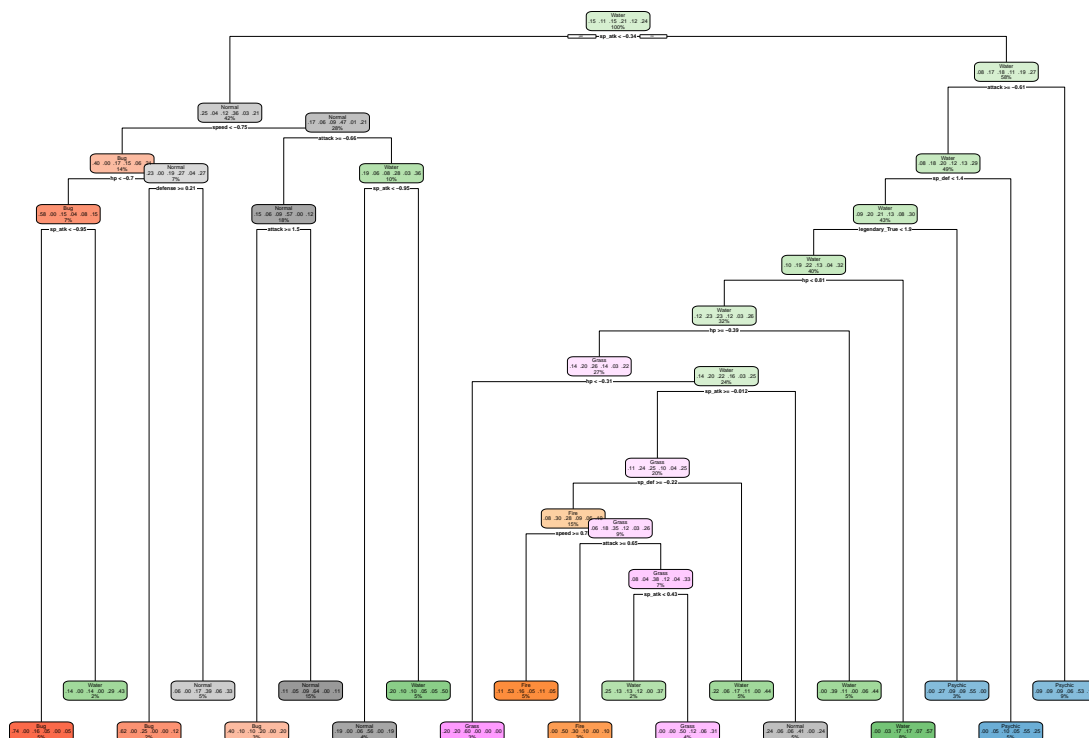
```
class_tree_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot()
```

```
## Warning: Cannot retrieve the data used to build the model (so cannot determine roundint and is.binary)
```

```
## To silence this warning:
```

```
##   Call rpart.plot with roundint=FALSE,
##   or rebuild the rpart model with model=TRUE.
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```

Exercise 5

Now set up a random forest model and workflow. Use the **ranger** engine and set **importance = "impurity"**. Tune **mtry**, **trees**, and **min_n**. Using the documentation for **rand_forest()**, explain in your own words what each of these hyperparameters represent.

Create a regular grid with 8 levels each. You can choose plausible ranges for each hyperparameter. Note that **mtry** should not be smaller than 1 or larger than 8. **Explain why not. What type of model would **mtry = 8** represent?**

```
rf_spec <- rand_forest(
  mtry = tune(),
  trees = tune(),
  min_n = tune()) %>%
  set_engine("ranger", importance = 'impurity') %>%
  set_mode('classification')

class_forest_rf <- workflow()%>%
  add_model(rf_spec)%>%
  add_recipe(pokemon_recipe)

pgram_grid<- grid_regular(mtry(range= c(1,8)),
  trees(range = c(100,1000)),
  min_n(range = c(1,10)),
  levels = 8)
```

- Explain in your own words what each of these hyperparameters represent.
- Answer:

- `mtry`: The number of predictors that will be randomly sampled at each split when we are creating the tree models.
- `trees`: The number of trees contained in the ensemble.
- `min_n`: The minimum number of data points in a node that are required for the node to be split further.
- Note that `mtry` should not be smaller than 1 or larger than 8. Explain why not.
- Answer: Notice that `mtry` is the number of predictors that will be randomly sampled at each split when we are creating the tree models. Since we only have 8 predictors, then we cannot make `mtry` greater than 8 or smaller than 1. According to the professor, if we do that, it won't be meaningful.
- What type of model would `mtry = 8` represent??
- Answer: `mtry = 8` represents the bagging model.

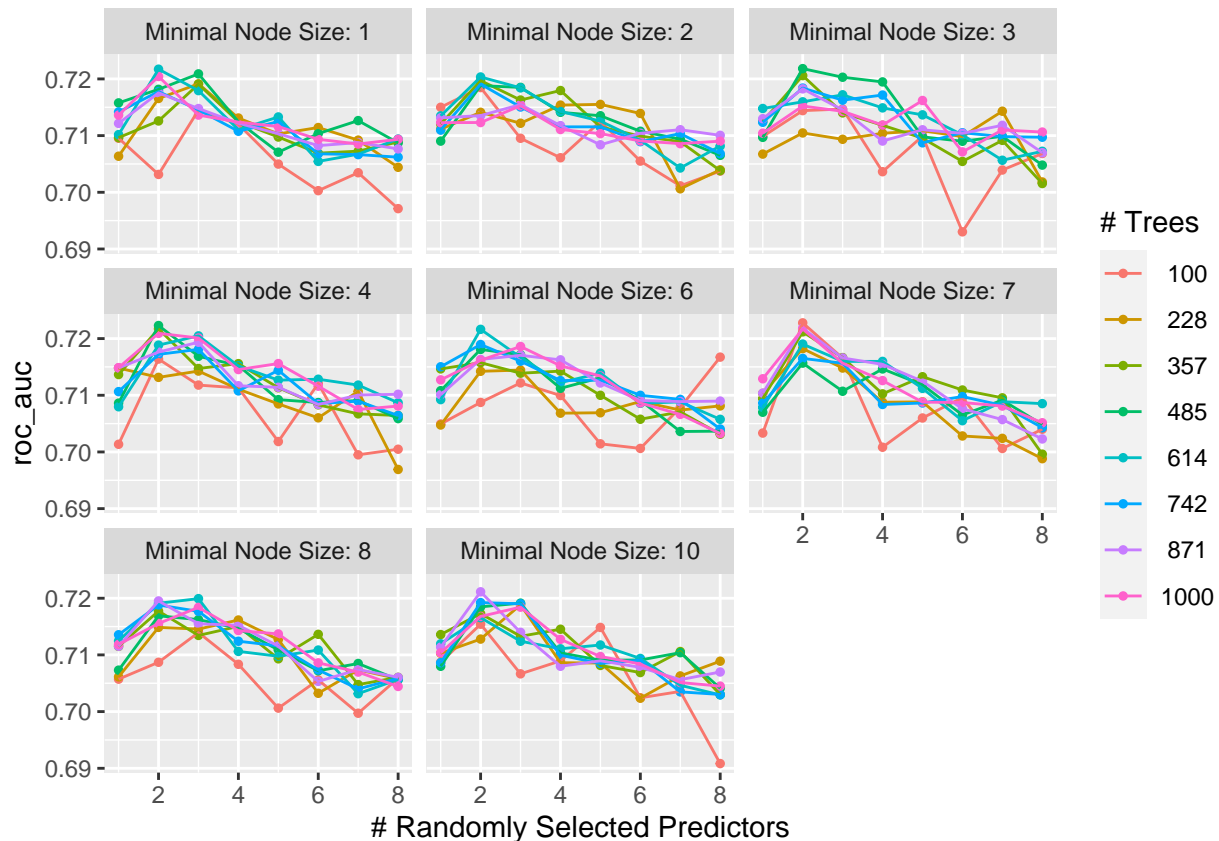
Exercise 6

Specify `roc_auc` as a metric. Tune the model and print an `autoplot()` of the results. What do you observe? What values of the hyperparameters seem to yield the best performance?

```
# Specify that the metric we want to optimize is `roc_auc`
forest_tune_res <- tune_grid(
  class_forest_rf,
  resamples = pokemon_folds,
  grid = pgram_grid,
  metrics = metric_set(roc_auc)
)
```

```
## ! Fold1: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold2: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold3: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold4: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold5: preprocessor 1/1: The following variables are not factor vectors and wil...

# Print an `autoplot()` of the results.
autoplot(forest_tune_res)
```



- What do you observe?
- Answer: According to the graph, if we fix the value of `min_n` (which is the minimal node size), we can see that as the number of randomly selected predictors (variable `mtry`) increases, the trend of `roc_auc` values of most of these models will finally decrease.
- What values of the hyperparameters seem to yield the best performance?
- Answer: According to the graph, we saw that when the number of randomly selected predictors (variable `mtry`) = 2, the number of tree (variable `trees`) = 100, and the minimal node size (variable `min_n`) = 7, we seem to yield the best performance.

Exercise 7

What is the `roc_auc` of your best-performing random forest model on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```
# What is the `roc_auc` of your best-performing boosted tree model on the folds?
arrange(collect_metrics(forest_tune_res), desc(mean))
```

```
## # A tibble: 512 x 9
##   mtry trees min_n .metric .estimator mean      n std_err .config
##   <int> <int> <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>
## 1     2   100     7 roc_auc hand_till 0.723     5 0.00925 Preprocessor1_Model~
## 2     2   485     4 roc_auc hand_till 0.722     5 0.00647 Preprocessor1_Model~
## 3     2   485     3 roc_auc hand_till 0.722     5 0.0112  Preprocessor1_Model~
## 4     2   871     7 roc_auc hand_till 0.722     5 0.00943 Preprocessor1_Model~
## 5     2   357     4 roc_auc hand_till 0.722     5 0.00968 Preprocessor1_Model~
## 6     2   614     1 roc_auc hand_till 0.722     5 0.0102  Preprocessor1_Model~
## 7     2  1000     7 roc_auc hand_till 0.722     5 0.0102  Preprocessor1_Model~
```

```
## 8      2    614      6 roc_auc hand_till  0.722      5 0.0109 Preprocessor1_Model~
## 9      2    357      7 roc_auc hand_till  0.721      5 0.00954 Preprocessor1_Model~
## 10     2    871     10 roc_auc hand_till  0.721      5 0.00820 Preprocessor1_Model~
## # ... with 502 more rows
```

```
best_complexity<- select_best(forest_tune_res, metric= 'roc_auc')
```

```
class_forest_final <- finalize_workflow(class_forest_rf, best_complexity)
```

- What is the roc_auc of your best-performing random forest model on the folds?
- Answer: The roc_auc of your best-performing random forest model on the folds is 0.723.

Exercise 8

Create a variable importance plot, using `vip()`, with your best-performing random forest model fit on the *training* set.


Which variables were most useful? Which were least useful? Are these results what you expected, or not?

```
class_forest_final_fit <- fit(class_forest_final, data = pokemon_train)
```

```
## Warning: The following variables are not factor vectors and will be ignored:
```

```
## `generation`
```

```
class_forest_final_fit%>%
  extract_fit_engine() %>%
  vip()
```



- Which variables were most useful?
- Answer: sp_atk is the most useful variable.

- Which were least useful?
- Answer: `legendary_True` (or `legendary`) is the least useful variable.
- Are these results what you expected, or not?
- Answer: Yes, these results are what I expected. Anyone who has watched Pokemon videos or played Pokemon games knows that `sp_atk` is very important to a Pokemon's strength, and the strength of a Pokemon has nothing to do with whether he/she is a legendary Pokemon.

Exercise 9

Finally, set up a boosted tree model and workflow. Use the `xgboost` engine. Tune `trees`. Create a regular grid with 10 levels; let `trees` range from 10 to 2000. Specify `roc_auc` and again print an `autoplot()` of the results.

What do you observe?

What is the `roc_auc` of your best-performing boosted tree model on the folds? *Hint: Use `collect_metrics()` and `arrange()`.*

```
boost_spec <- boost_tree(trees = tune(), tree_depth = 4) %>%
  set_engine("xgboost") %>%
  set_mode("classification")


boost_rf <- workflow()%>%
  add_model(boost_spec)%>%
  add_recipe(pokemon_recipe)

pgram3_grid<- grid_regular(trees(range = c(10,2000) ),
                           levels = 10)

# Specify that the metric we want to optimize is `roc_auc`
boost_tune_res <- tune_grid(
  boost_rf,
  resamples = pokemon_folds,
  grid = pgram3_grid,
  metrics = metric_set(roc_auc)
)
```

```
## ! Fold1: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold2: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold3: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold4: preprocessor 1/1: The following variables are not factor vectors and wil...
## ! Fold5: preprocessor 1/1: The following variables are not factor vectors and wil...

# Print an `autoplot()` of the results.
autoplot(boost_tune_res)
```



What is the `roc_auc` of your best-performing boosted tree model on the folds?
`arrange(collect_metrics(boost_tune_res), desc(mean))`

```
## # A tibble: 10 x 7
##   trees .metric .estimator  mean     n std_err .config
##   <int> <chr>    <chr>    <dbl> <int>  <dbl> <chr>
## 1  2000 roc_auc hand_till  0.698     5 0.00846 Preprocessor1_Model10
## 2   231 roc_auc hand_till  0.698     5 0.00700 Preprocessor1_Model02
## 3  1778 roc_auc hand_till  0.697     5 0.00864 Preprocessor1_Model09
## 4  1557 roc_auc hand_till  0.697     5 0.00838 Preprocessor1_Model08
## 5  1336 roc_auc hand_till  0.697     5 0.00813 Preprocessor1_Model07
## 6   673 roc_auc hand_till  0.696     5 0.00779 Preprocessor1_Model04
## 7   452 roc_auc hand_till  0.696     5 0.00708 Preprocessor1_Model03
## 8  1115 roc_auc hand_till  0.696     5 0.00789 Preprocessor1_Model06
## 9   894 roc_auc hand_till  0.695     5 0.00742 Preprocessor1_Model05
## 10    10 roc_auc hand_till  0.686     5 0.0171  Preprocessor1_Model01
```

- What do you observe?
- Answer: Based on the graph we get, we can see that:
 - when the number of trees is less than 231, as the number of trees increases, the value of roc_auc will increase;
 - when the number of trees is greater than 231 but less than 452, as the number of trees increases, the value of roc_auc will decrease;
 - when the number of trees is greater than 452 but less than 894, as the number of trees increases, the value of roc_auc will increase a little first and then decrease;

- when the number of trees is greater than 894, as the number of trees increases, the value of `roc_auc` will increase.
- What is the `roc_auc` of your best-performing boosted tree model on the folds?
- Answer: The `roc_auc` of your best-performing boosted tree model on the folds is 0.698.

Exercise 10

Display a table of the three ROC AUC values for your best-performing pruned tree, random forest, and boosted tree models. Which performed best on the folds? Select the best of the three and use `select_best()`, `finalize_workflow()`, and `fit()` to fit it to the *testing* set.

Print the AUC value of your best-performing model on the testing set. Print the ROC curves. Finally, create and visualize a confusion matrix heat map.

Which classes was your model most accurate at predicting? Which was it worst at?

```
# Display a table of the three ROC AUC values for your
# best-performing pruned tree, random forest, and boosted tree models.
value <- c(arrange(collect_metrics(tune_res), desc(mean))[1,4],
           arrange(collect_metrics(forest_tune_res), desc(mean))[1,6],
           arrange(collect_metrics(boost_tune_res), desc(mean))[1,4])
cnames <- c('ROC AUC values of pruned tree', 'ROC AUC values of random forest',
            'ROC AUC values of boosted tree')
rnames <- 'values'
table<- matrix(value, nrow = 1, ncol = 3, byrow = TRUE, dimnames = list(rnames,cnames))
table

##          ROC AUC values of pruned tree ROC AUC values of random forest
## values 0.6437168                      0.7227707
##          ROC AUC values of boosted tree
## values 0.698091

# Which performed best on the folds?
# random forest model performed best on the folds
# Select the best of the three and use `select_best()`, `finalize_workflow()`,
# and `fit()` to fit it to the *training* set.
best_complexity<- select_best(forest_tune_res, metric= 'roc_auc')

class_forest_final2 <- finalize_workflow(class_forest_rf, best_complexity)

class_forest_final_fit2 <- fit(class_forest_final, data = pokemon_train)

## Warning: The following variables are not factor vectors and will be ignored:
## `generation`

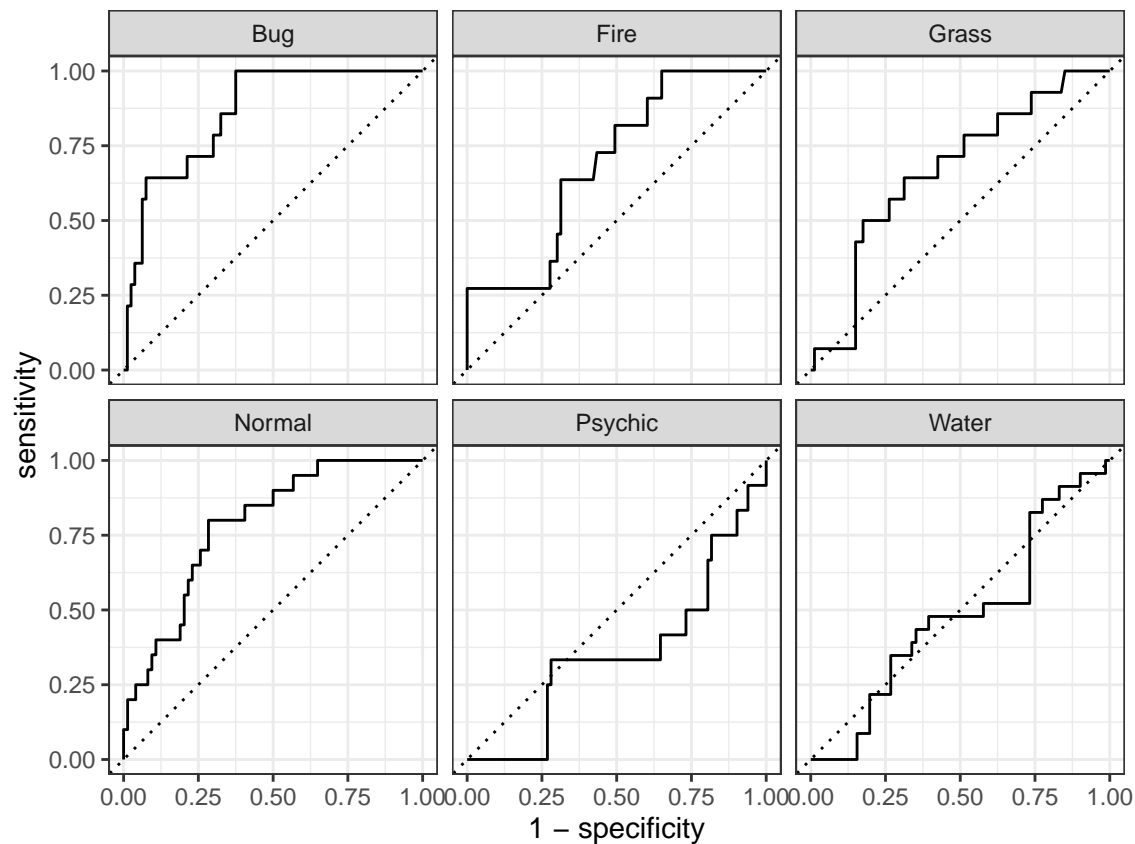
# Print the AUC value of your best-performing model on the testing set.

augment(class_forest_final_fit2, new_data = pokemon_test) %>%
  roc_auc(type_1, .pred_Bug, .pred_Fire, .pred_Grass, .pred_Normal, .pred_Water, .pred_Psychic)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc hand_till    0.639
```

```
# Print the ROC curves.
```

```
augment(class_forest_final_fit2, new_data = pokemon_test) %>% roc_curve(type_1, .pred_Bug, .pred_Fire,
```



```
# Finally, create and visualize a confusion matrix heat map.
```

```
augment(class_forest_final_fit2, new_data = pokemon_test) %>%  
  conf_mat(truth = type_1, estimate = .pred_class)%>%  
  autoplot("heatmap")
```


Prediction	Bug -	6	0	1	4	1	0
	Fire -	0	3	2	1	0	1
	Grass -	0	1	2	0	0	2
	Normal -	7	1	0	9	3	4
	Psychic -	0	3	2	2	7	2
	Water -	1	3	7	4	1	14
		Bug	Fire	Grass	Normal	Psychic	Water
		Truth					

Print the AUC value of your best-performing model on the testing set. * Answer:0.639

- Which classes was your model most accurate at predicting? Which was it worst at?
- Answer:
- In order to answer this question, we need to calculate $\frac{\# \text{ the number of a specific class which be predicted correctly}}{\# \text{ the total number of predictions}}$.
- Bug: $\frac{6}{(6 + 0 + 1 + 4 + 1 + 0)} = 0.5$
- Fire: $\frac{3}{(0 + 3 + 2 + 1 + 0 + 1)} = 0.4286$
- Grass: $\frac{2}{(0 + 1 + 2 + 0 + 0 + 2)} = 0.4$
- Normal: $\frac{9}{(7 + 1 + 0 + 9 + 3 + 4)} = 0.375$
- Psychic: $\frac{7}{(0 + 3 + 2 + 2 + 7 + 2)} = 0.4375$
- Water: $\frac{14}{(1 + 3 + 7 + 4 + 1 + 14)} = 0.46667$
- Based on the calculations, we know that my model most accurate at predicting at Class Bug, and my model worst accurate at predicting at Class Normal.